

Real-Time Systems

Systèmes Temps-Réel embarqués



Real-time embedded systems



Embedded systems:

- « An application whose prime function is not that of information processing, but which nevertheless requires information processing in order to carry out their prime function » – Real-Time Systems and programming languages, A. Burns and A. Wellings
- Ex : Microprocessor-controlled washing machine

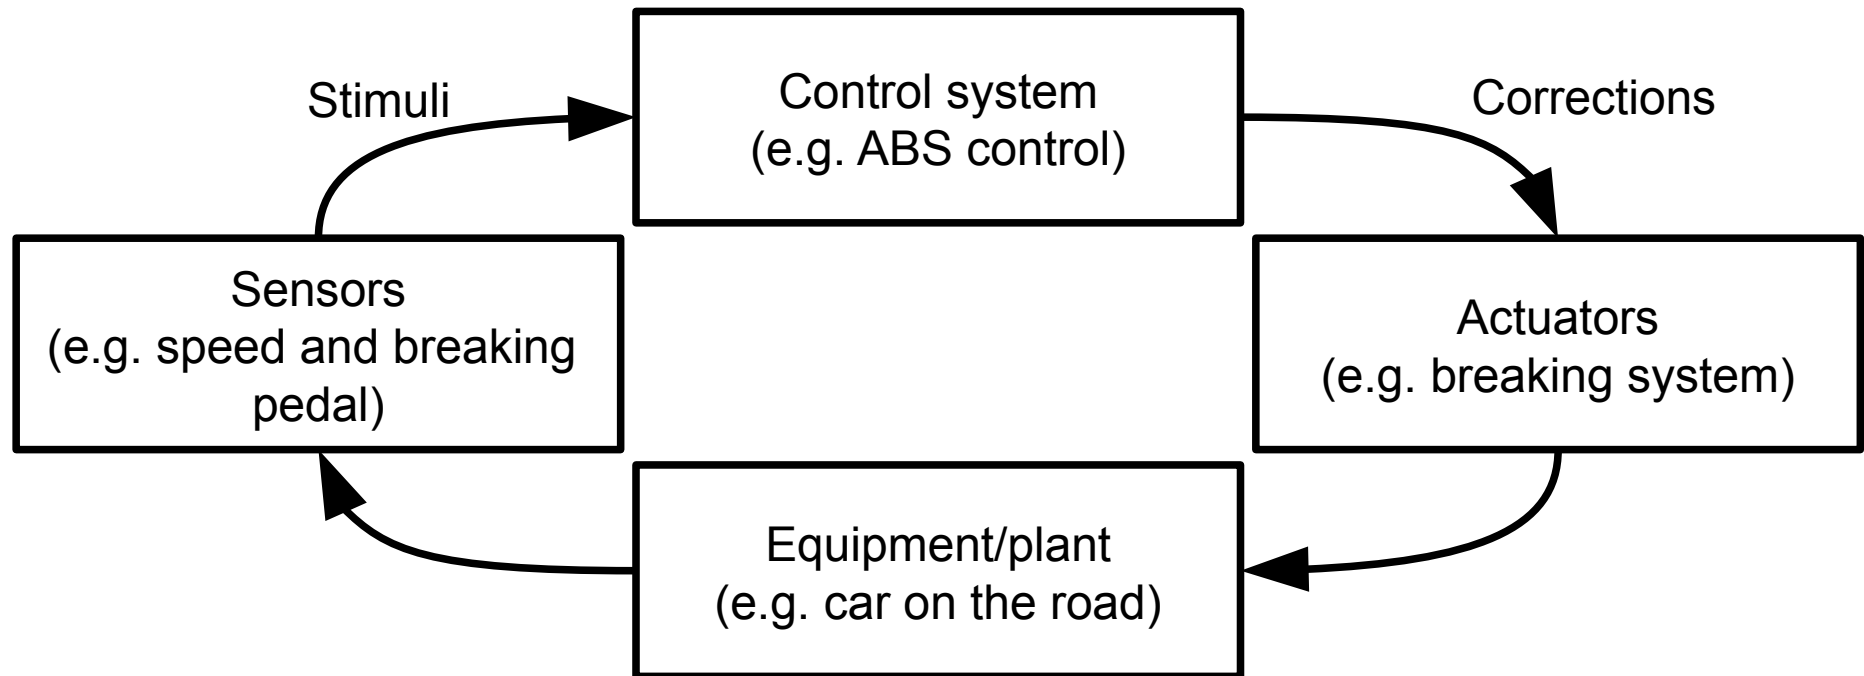
Real-time systems:

- « Any system in which the time at which output is produced is significant » – Oxford dictionary of computing
- Outputs must be produced within a certain duration (**deadline**) after input stimuli have been produced by the **environment** (subcategory of **reactive systems**)
- Real-time system = functional correctness + **temporal correctness**
- Ex: Control system for airbags, ABS, etc.

Real-time embedded systems



Ex: embedded control systems:



Real-time embedded systems



Soft Real-time:

- Deadlines can be missed occasionally, or service can occasionally be delivered late
- Ex: telemetry report system, car window controller, multimedia application, etc.

Hard Real-Time:

- Any deadline miss is considered a system failure and may be **catastrophic**
- Critical systems : systems for which a timing failure may cause human harm or loss, environmental disaster or expensive mission failure
- Ex: aircraft control system, nuclear plant, etc.

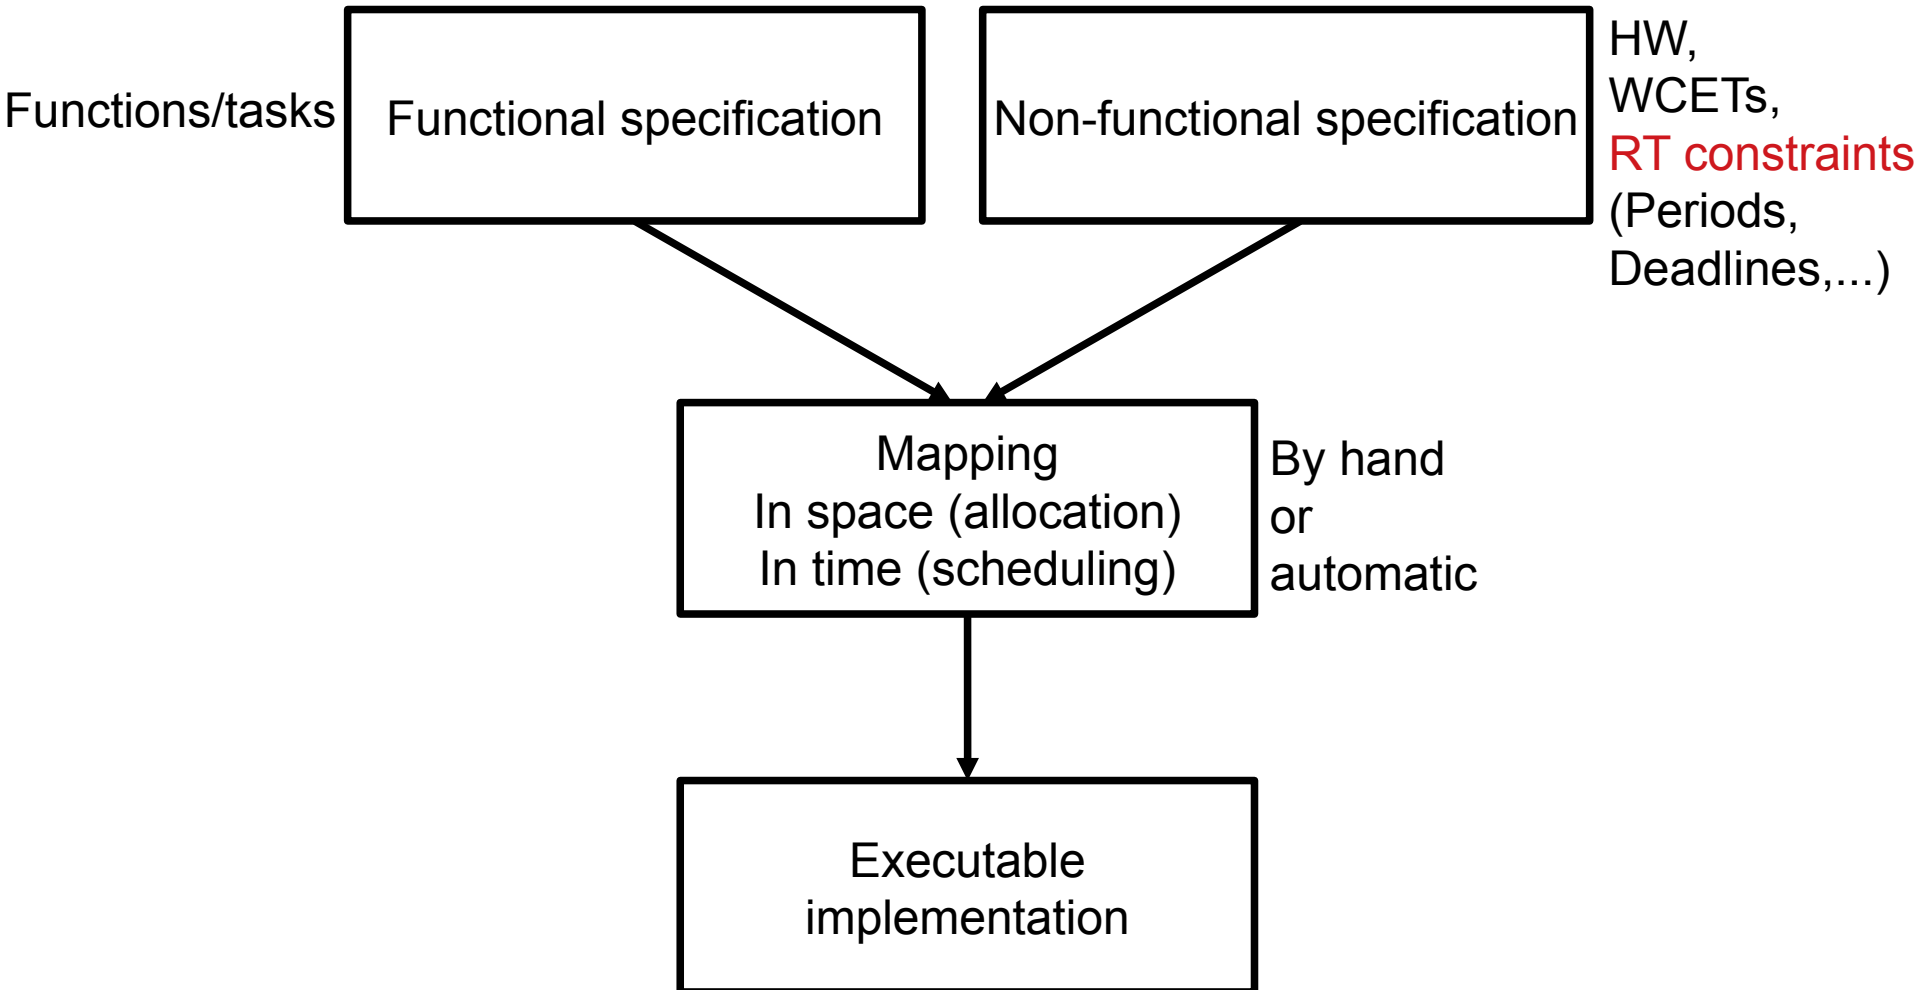
Real-time embedded systems



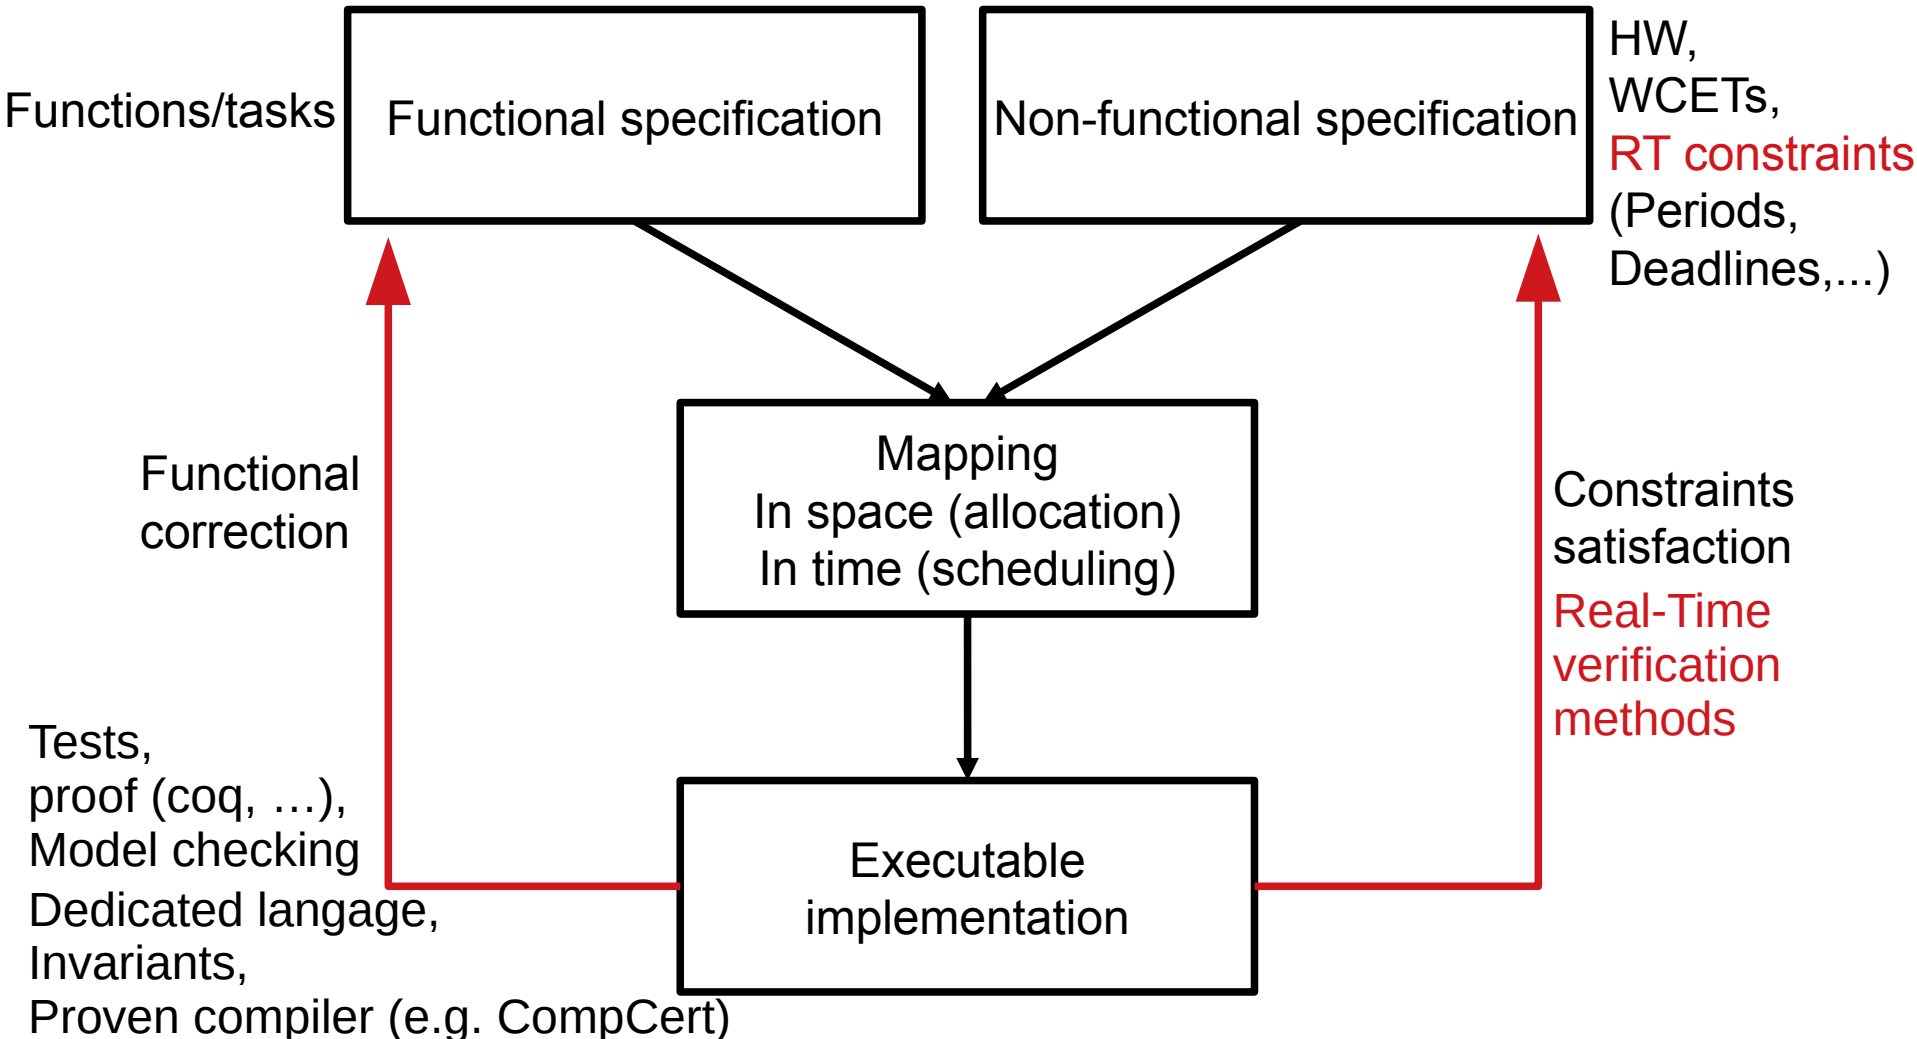
Engineering a real-time system:

- Specification: functional and non-functional requirements
 - Functional: What the system must compute ? Defining **tasks**
 - Non-functional: On what hardware ? With which temporal requirements ?
- Implementation: producing the system with constraints
 - Languages: non-synchronous (C, ADA, ...) vs synchronous (lustre/SCADE, esterel, ...)
 - Cyclic executive vs concurrent. Linking **specification tasks** to **execution threads/processes**
 - Generating static code vs generating scheduler configuration
- Verification: verifying the requirements on the produced system
 - Functional (not part of this course)
 - Non-functional (**schedulability analysis**)

Real-time embedded systems



Real-time embedded systems



Specification



Real-Time requirements sources

- Laws of nature:
 - Objects in motion, inertia, etc.
 - => End-to-end response time constraints/deadlines
- Control theory:
 - Execution frequency/period
- Hardware limitations:
 - Sensor sampling rate
 - Actuation rate
 - => Periods

Specification



Our running example: a simplified car control system

- Functional requirements:
 - Implement an Engine Control Unit:
 - Samples inputs from a pedal angle sensor, a speed sensor and an engine rotation sensor
 - Writes to a controller in the engine to control the rotation speed
 - Implement an Airbag Control Unit:
 - Samples inputs from a collision detector
 - Triggers the airbag if necessary
- Real-time requirements:
 - Sampling rate for the sensors:
 - Pedal angle, speed: 10ms
 - Engine rotation: 20ms
 - Collision detection: 60ms
 - Actuation rates:
 - ECU function (control theory): 40ms
 - Airbag function (bodies in motion): 60ms

Specification



Most common abstraction: real-time tasks

- Tasks: the functional units of the system (e.g. functions or groups of functions)
- Periodic tasks: mostly used
 - Period (T_i): the duration between two consecutive activations of a task
 - Relative deadline (a.k.a. deadline) (D_i): the duration between a task instance arrival date and the date at which it must complete its execution
 - Execution time budget (C_i) => Worst-case execution time
- Periodicity induces multiple instances (jobs)
 - Arrival date (A_{ij}): the date at which job j of task i is ready for scheduling (a.k.a. release date)
 - Absolute deadline: the date at which a job must complete its execution

Specification



Definitions:

- Tasks: the functional units of the system (e.g. functions or groups of functions)
- Periodic tasks: mostly used
 - Sometimes characterized only with T_i and C_i (when $D_i=T_i$)

Specification



Definitions:


- Tasks: the functional units of the system (e.g. functions or groups of functions)
- Periodic tasks: mostly used
- Sporadic tasks:
 - Not periodic, but minimum inter-arrival duration is known
 - E.g. fault-tolerance routines, some sensor polling tasks
- Aperiodic tasks:
 - Can happen anytime, makes verification hard
 - Can be tackled by adding periodic « server » tasks

Specification



Example:

- $\tau_0 : T_0=10, D_0=6, C_0=3$



Specification



Our running example: a simplified car control system

■ Functional requirements:

- Implement an Engine Control Unit:
 - Samples inputs from a pedal angle sensor, a speed sensor and an engine rotation sensor
 - Writes to a controller in the engine
- Implement an Airbag Control Unit:
 - Samples inputs from a collision detector
 - Triggers the airbag if necessary

■ Real-time requirements:

- Sampling rate for the sensors:
 - Pedal angle, speed: 10ms
 - Engine rotation: 20ms
 - Collision detection: 60ms
- Actuation rates:
 - ECU function (control theory): 40ms
 - Airbag function (bodies in motion): 60ms

How to capture the requirements in the periodic task model ?

Specification



Our running example: a simplified car control system

■ Real-time requirements:

- Sampling rate for the sensors:

- Pedal angle, speed: 10ms
- Engine rotation: 20ms
- Collision detection: 60ms

- Actuation rates:

- ECU function (control theory): 30ms
- Airbag function (bodies in motion): 60ms

Engine Control
Unit
(ECU)

$T_0=10\text{ms}$?

$T_0=30\text{ms}$?

Airbag Control
Unit

$T_1=60\text{ms}$

■ Original model: One task per functionality (coarse-grain tasks)

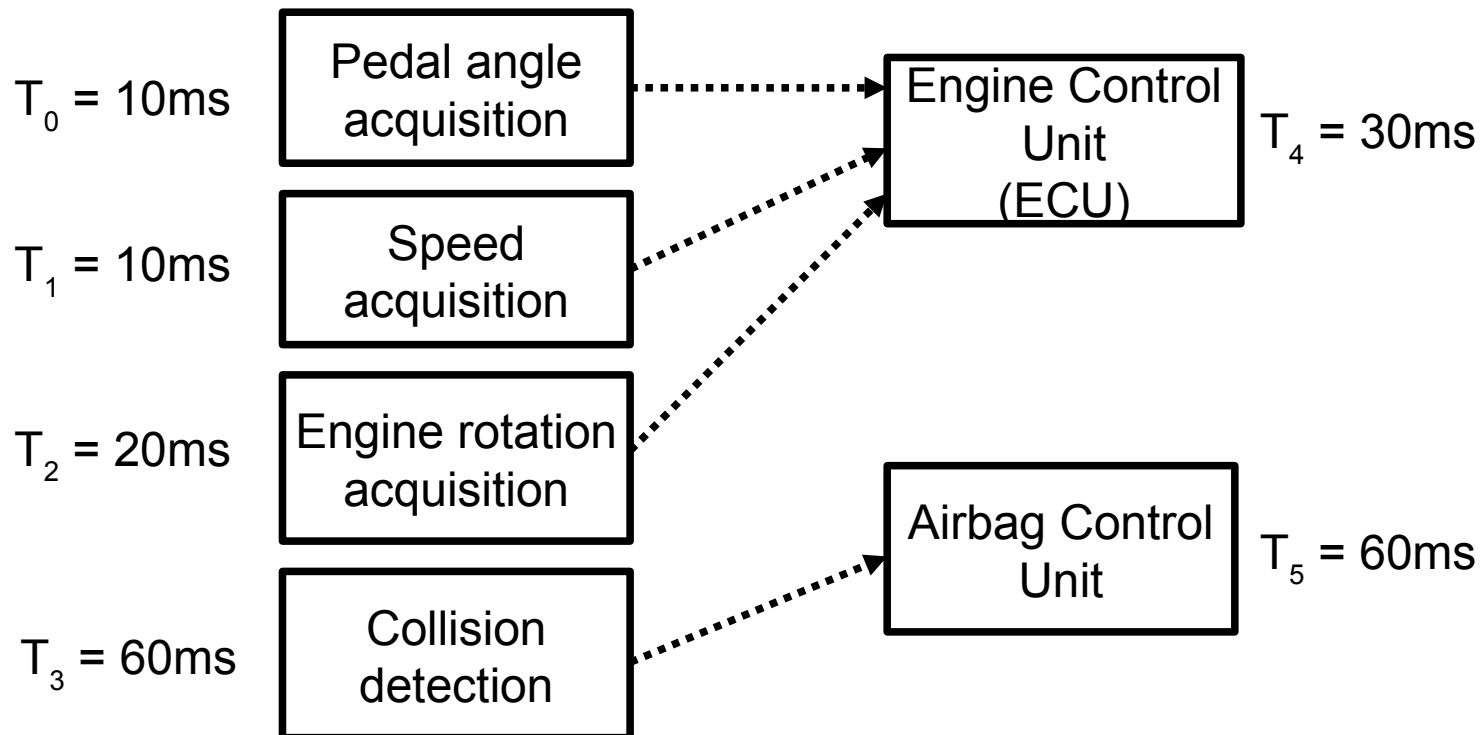
- Each is responsible for input acquisition, correction computation and actuation
- Very simple system design, complexity hidden in the tasks
 - In particular, different sampling rates
 - => Tasks may be harder/sometimes impossible to implement
 - => Tasks WCET are harder to derive
- => Strong abstraction, usually leads to pessimism overhead in validation phase
- => Hard to optimize at system-level

Specification



Our running example: a simplified car control system

- Model relaxation: one task per function (fine grain tasks)
 - For now, we do not take into account inter-task communications



Specification



Our running example: a simplified car control system

- Model relaxation: one task per function (fine grain tasks)
 - For now, we do not take into account inter-task communications
 - System design is more complex, tasks are simpler to implement
 - Abstraction is weaker
 - => system-level analysis is harder, but optims are possible
 - => small WCET overhead when analyzing tasks
- Choosing the task system:
 - Transforms requirements into constraints
 - Constraints may be stricter than requirements (cost of abstraction): cf original model in the example
 - Is the first step of implementation (limit between specification and implementation is blurred)

Implementation



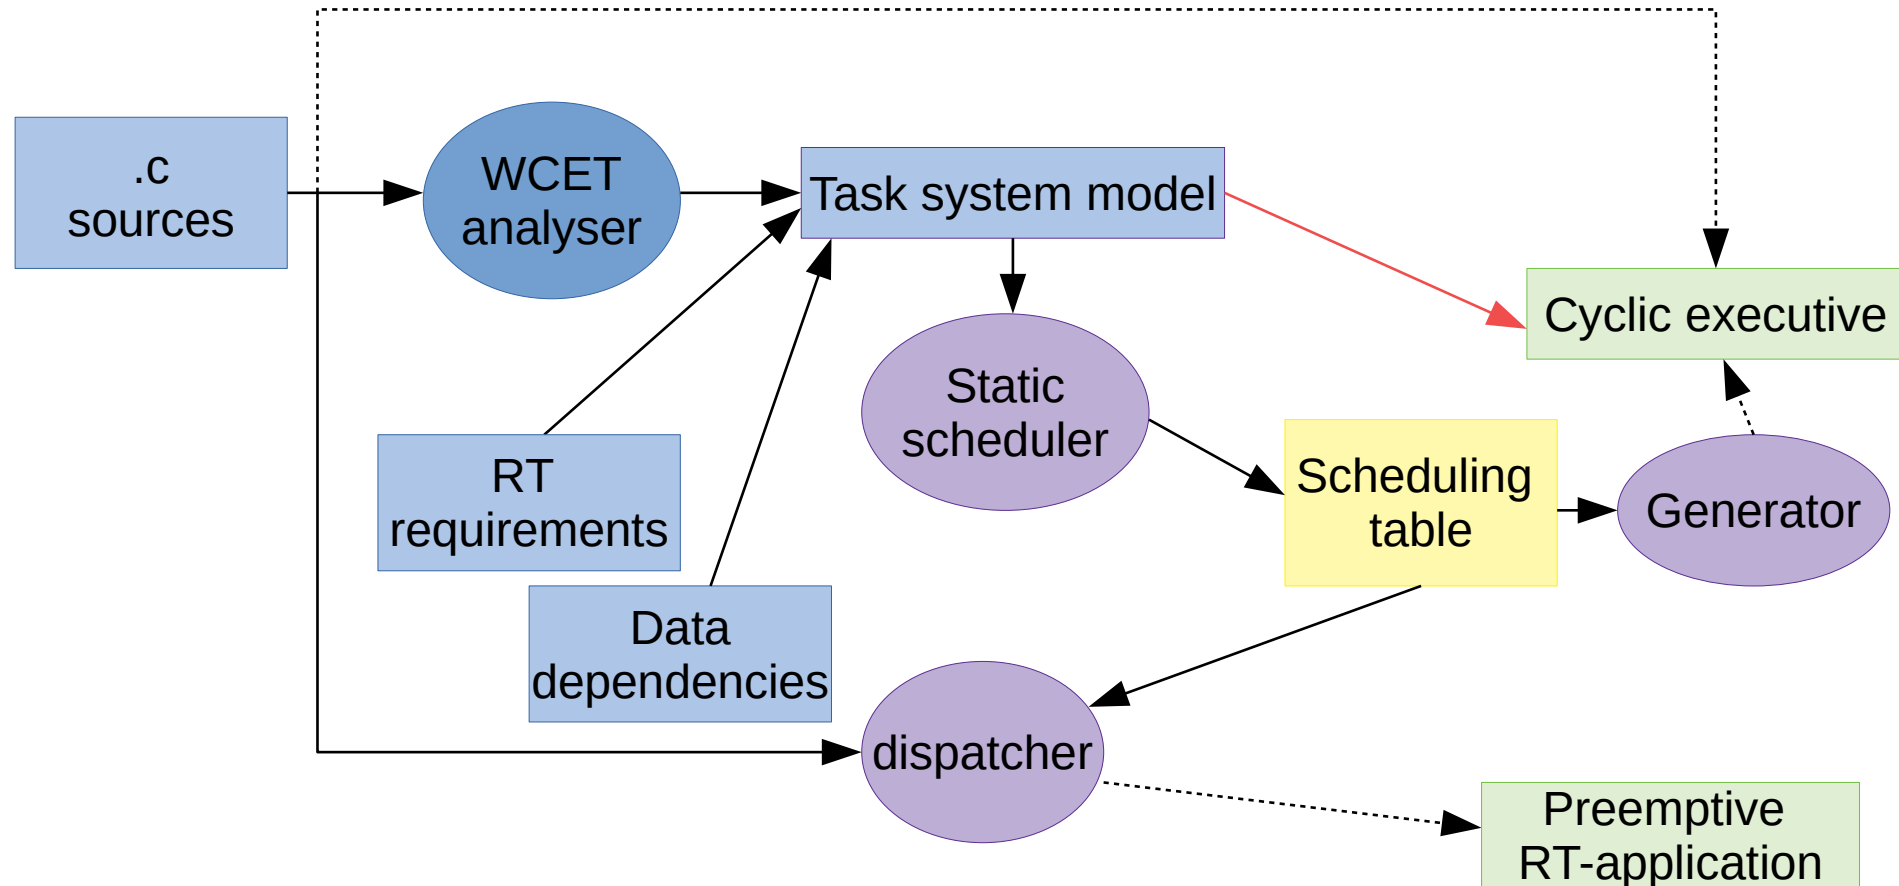
Two main implementation families:

- Static implementation: everything is fixed offline
 - Produce a static schedule i.e. order of execution of the tasks, sometimes also their starting dates
 - Produce a cyclic executive, then check it
- OR
- Produce a scheduling table with the start date of each task and the code for each task (processes or threads)
- Easy to verify the RT properties => good for very critical parts
- Dynamic implementation: decisions are made online (during the execution)
 - Implement/choose a scheduler to make the decisions at runtime
 - Choose a scheduling policy
 - Produce tasks code (processes or threads)
 - Verification is harder but implementation is easier, more modular and robust

Static (off-line) scheduling



Static scheduling:



Static (off-line) scheduling



Cyclic executive:

- Basically, one infinite loop which contains calls to the tasks' functions
 - Only one process/thread

```
Loop:                                //Repeat indefinitely
    acquire_inputs();                //Sensor polling
    compute_corrections();
    write_outputs();                 //Update actuators
```

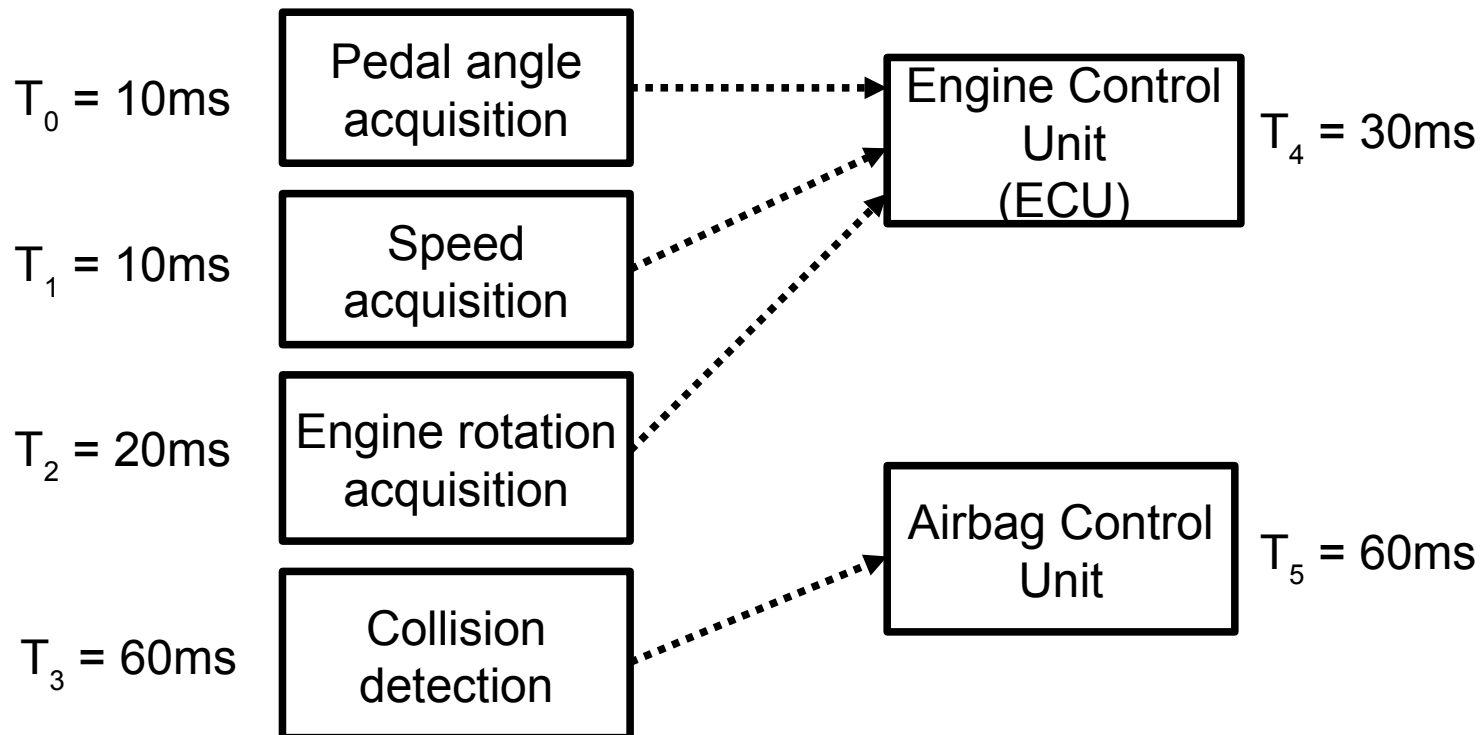
- The main loop gives the activation period for the whole system
 - In the absence of synchronizations, the sum of the WCETs of the tasks' functions gives the response time AND the worst case period

Static (off-line) scheduling



Our running example: a simplified car control system

- Second version: one task per function (fine grain tasks)
 - For now, inter-task communications can be ignored (using buffers)



Static (off-line) scheduling



Cyclic executive:

loop :

```
    pedal_angle();  
    speed();  
    if(i%2==0) engine_rotation();  
    if(i%6==0) collision_detection();  
    if(i%3==0) ECU();  
    if(i%6==0) airbag();  
    i++;  
    wait_10ms(); // wait for timer interrupt  
                // or busy wait until 10ms is reached
```

Static (off-line) scheduling



Cyclic executive:

loop :

```
    pedal_angle();
```

```
    speed();
```

```
    if(i%2==0) engine_rotation();
```

```
    if(i%6==0) collision_detection();
```

```
    if(i%3==0) ECU();
```


```
    if(i%6==0) airbag();
```

```
    i++;
```

```
    wait_10ms(); // wait for timer interrupt
```

```
                // or busy wait until 10ms is reached
```

Period of loop is the period of the fastest task



Static (off-line) scheduling



Cyclic executive:

- Simple verification method: sum of worst-case execution times must be less than the period of the loop (and/or less than the smallest deadline)
 - Regardless of the order of function calls in the loop
 - Sufficient criteria, but not necessary
 - Exact criteria when all tasks have the same period

$$\forall i, \sum_j C_j \leq D_i$$

$$\forall i, \sum_j C_j \leq T_i$$

- Exact verification method: build the worst-case schedule of the task system (one hyperperiod when all tasks start at date 0)
 - If no task exceeds its own deadline, the system is safe

Static (off-line) scheduling



Cyclic executive: sufficient vs exact method

loop :

$C_0=1$ `pedal_angle();`

$C_1=1$ `speed();`

$C_2=2$ `if(i%2==0) engine_rotation();`

$C_3=2$ `if(i%6==0) collision_detection();`


$C_4=3$ `if(i%3==0) ECU();`

$C_5=4$ `if(i%6==0) airbag();`

`i++;`

`wait_10ms(); // wait for timer interrupt`

`// or busy wait until 10ms is reached`



Static (off-line) scheduling



Cyclic executive: Problem with long tasks

loop :

$C_0=1$ `pedal_angle();`

$C_1=1$ `speed();`

$C_2=2$ `if(i%2==0) engine_rotation();`

$C_3=2$ `if(i%6==0) collision_detection();`

$C_4=3$ `if(i%3==0) ECU();`

$C_5=12$ `if(i%6==0) airbag();`


`i++;`

`wait_10ms(); // wait for timer interrupt`

`// or busy wait until 10ms is reached`

Period of loop is the period of the fastest task => if a task takes longer than this period we have a deadline miss

Deadline miss !!! => Schedule with larger period or add preemptions



Static (off-line) scheduling





Scheduling table:

- Represents the starting dates of the tasks in each hyperperiod (lcm of tasks periods)
 - Program object
 - Used to generate and/or verify a simple cyclic executive
 - Verify: create cyclic executive, then derive corresponding table
 - Generate: create table first, then corresponding executive
 - Also used as a parameter to a **dispatcher**
 - Can support preemptions
 - Starting dates and order of computation are fixed at compile time, no decision is made online by the dispatcher
-
- A scheduling table is its own **validation criteria**: if no task violates its deadline in the table, the implementation is safe

Static (off-line) scheduling



Scheduling table



Static (off-line) scheduling



Scheduling table:

- How to build the scheduling table ?
 - NP-hard problem
 - => for small task systems: by hand, using ILP or heuristics
 - => for large task systems: use heuristics
- When there is no inter-task dependence
 - Simulate an online scheduling policy for one hyperperiod
 - E.g. RM, DM, EDF, LLF, etc.
- When there are inter-task dependencies
 - Exhibit all jobs for one hyperperiod in your task model (including release dates and dependencies to other jobs)
 - Schedule ASAP while following the partial order dictated by the dependencies
=> List scheduling
 - In case of multiple candidate tasks, try all and choose according to a cost function (greedy algorithm)
 - Depending on the size of the task system, use backtracking

Static (off-line) scheduling



Scheduling table:

- Preemptive rate-monotonic scheduling
 - RM is initially an online scheduling algorithm
 - In RM, each task is assigned a **static (fixed) priority**
 - Priorities are based on tasks **periods**
 - The smaller the period, the higher the priority
 - $T_i < T_j \Rightarrow \text{priority}(\tau_i) > \text{priority}(\tau_j)$
 - Under certain hypothesis, RM is optimal (in the sense of RT scheduling)

Static (off-line) scheduling



Scheduling table:

- Preemptive rate-monotonic scheduling
 - Select tasks by increasing activation periods
 - Schedule each job of the task for the hyperperiod ASAP

$$T_0 = 10\text{ms}$$
$$C_0 = 1\text{ms}$$

Pedal angle
acquisition

$$T_1 = 10\text{ms}$$
$$C_1 = 1\text{ms}$$

Speed
acquisition

$$T_2 = 20\text{ms}$$
$$C_2 = 2\text{ms}$$

Engine rotation
acquisition

$$T_3 = 60\text{ms}$$
$$C_3 = 2\text{ms}$$

Collision
detection

Engine Control
Unit
(ECU)

$$T_4 = 30\text{ms}$$
$$C_4 = 3\text{ms}$$

Airbag Control
Unit

$$T_5 = 60\text{ms}$$
$$C_5 = 12\text{ms}$$


Static (off-line) scheduling



Scheduling table:

- Preemptive rate-monotonic scheduling
 - Select tasks by increasing activation periods
 - Schedule each job of the task for the hyperperiod ASAP

Start with an empty table of size hyperperiod



Static (off-line) scheduling




Scheduling table:

- Preemptive rate-monotonic scheduling
 - Select tasks by increasing activation periods
 - Schedule each job of the task for the hyperperiod ASAP

Select and schedule ASAP task with lowest period

$$T_0 = 10\text{ms}$$
$$C_0 = 1\text{ms}$$

Pedal angle
acquisition



Static (off-line) scheduling




Scheduling table:

- Preemptive rate-monotonic scheduling
 - Select tasks by increasing activation periods
 - Schedule each job of the task for the hyperperiod ASAP

Select and schedule ASAP task with next lowest period $T_1 = 10\text{ms}$
 $C_1 = 1\text{ms}$

Speed
acquisition



Static (off-line) scheduling




Scheduling table:

- Preemptive rate-monotonic scheduling
 - Select tasks by increasing activation periods
 - Schedule each job of the task for the hyperperiod ASAP

Select and schedule ASAP task with next lowest period

$$T_2 = 20\text{ms}$$
$$C_2 = 2\text{ms}$$

Engine rotation
acquisition



Static (off-line) scheduling




Scheduling table:

- Preemptive rate-monotonic scheduling
 - Select tasks by increasing activation periods
 - Schedule each job of the task for the hyperperiod ASAP

Select and schedule ASAP task with next lowest period

Engine Control
Unit
(ECU)

$T_4 = 30\text{ms}$
 $C_4 = 3\text{ms}$



Static (off-line) scheduling




Scheduling table:

- Preemptive rate-monotonic scheduling
 - Select tasks by increasing activation periods
 - Schedule each job of the task for the hyperperiod ASAP

Select and schedule ASAP task with next lowest period
(I could have selected the airbag control task instead
resulting in a different, yet valid, schedule)

$$T_3 = 60\text{ms}$$
$$C_3 = 2\text{ms}$$

Collision
detection



Static (off-line) scheduling




Scheduling table:

- Preemptive rate-monotonic scheduling
 - Select tasks by increasing activation periods
 - Schedule each job of the task for the hyperperiod ASAP

Select and schedule ASAP task with next lowest period

Airbag Control Unit

$T_5 = 60\text{ms}$
 $C_5 = 12\text{ms}$



Static (off-line) scheduling



Scheduling table:

■ Preemptive rate-monotonic scheduling

- Select tasks by increasing activation periods
- Schedule each job of the task for the hyperperiod ASAP
- **Sufficient** schedulability criterion

$$U = \sum_{i=1}^n \frac{C_i}{T_i} \leq n(2^{1/n} - 1)$$

- $n \rightarrow +\infty : U \leq \ln 2 \approx 0.69$
- Only if :
 - Preemptive
 - No shared lock (no critical section)
 - $T_i = D_i$ for all i
 - All preemption costs are accounted for in the C_i (CRPD and context switch)
- If all periods are harmonic, $U \leq 1$ is the sufficient criterion

Static (off-line) scheduling



Scheduling table:

- Preemptive rate-monotonic scheduling
 - Select tasks by increasing activation periods
 - Schedule each job of the task for the hyperperiod ASAP
 - **Exact** schedulability criterion: worst-case response-time analysis
 - For each task, compute response-time of the first generated job
 - If all tasks start at date 0, the worst-case processor demand happens at date 0 (a.k.a synchronous release) => if all jobs generated at date 0 finish before their deadline, then any job will finish before its deadline => the system is schedulable

Static (off-line) scheduling



Scheduling table:

- Preemptive rate-monotonic scheduling
 - Select tasks by increasing activation periods
 - Schedule each job of the task for the hyperperiod ASAP
 - **Exact** schedulability criterion: worst-case response-time analysis
 - How to compute the WCRT of each job ?
 - A job executes as soon as it is released, but can be delayed by tasks with higher or same priority (all of them in the worst case).
 - The WCRT is thus equal to the job's execution budget (WCET) plus the interference coming from higher than or equal priority tasks.

Static (off-line) scheduling



Scheduling table:

- Preemptive rate-monotonic scheduling
 - Select tasks by increasing activation periods
 - Schedule each job of the task for the hyperperiod ASAP
 - **Exact** schedulability criterion: worst-case response-time analysis
 - $\tau_0 : T_0=10, C_0=3$
 - In isolation, WCRT $R_0 = C_0$




Static (off-line) scheduling



Scheduling table:

- Preemptive rate-monotonic scheduling
 - Select tasks by increasing activation periods
 - Schedule each job of the task for the hyperperiod ASAP
 - **Exact** schedulability criterion: worst-case response-time analysis
 - $\tau_0 : T_0=10, C_0=3$
 - $\tau_1 : T_1=5, C_1=1$
 - $R_1 = C_1 = 1$
 - $R_0 = C_0 + C_1 = 4$



Static (off-line) scheduling




Scheduling table:

■ Preemptive rate-monotonic scheduling

- Select tasks by increasing activation periods
- Schedule each job of the task for the hyperperiod ASAP
- **Exact** schedulability criterion: worst-case response-time analysis
 - $\tau_0 : T_0=10, C_0=3$
 - $\tau_1 : T_1=5, C_1=1$
 - $\tau_2 : T_2=5, C_2=2$

$$\text{sum}(C_i/T_i)=9/10$$
$$3 \cdot (2^{1/3}-1) \sim 0,77$$

Sufficient condition does not hold



Static (off-line) scheduling




Scheduling table:

- Preemptive rate-monotonic scheduling
 - Select tasks by increasing activation periods
 - Schedule each job of the task for the hyperperiod ASAP
 - **Exact** schedulability criterion: worst-case response-time analysis
 - $j > i \iff T_j \leq T_i$

$$R_i = C_i + \sum_{j=i+1}^n \left\lceil \frac{R_i}{T_j} \right\rceil * C_j$$

How many jobs from τ_j arrive in an interval of duration R_i ?



Static (off-line) scheduling



Scheduling table:

■ Preemptive rate-monotonic scheduling

- Select tasks by increasing activation periods
- Schedule each job of the task for the hyperperiod ASAP
- **Exact** schedulability criterion: worst-case response-time analysis
 - $j > i \iff T_j \leq T_i$

$$R_i = C_i + \sum_{j=i+1}^n \left\lceil \frac{R_i}{T_j} \right\rceil * C_j$$

How many jobs from τ_j arrive in an interval of duration R_i ?

- Must be computed as a fixed-point !

Static (off-line) scheduling



Scheduling table:

■ Preemptive rate-monotonic scheduling


- Select tasks by increasing activation periods
- Schedule each job of the task for the hyperperiod ASAP
- **Exact** schedulability criterion: worst-case response-time analysis

- $j > i \iff T_j \leq T_i$

- Fixed-point computation (step 1):

- $R_0^1 = C_0 = 3$

$$R_i = C_i + \sum_{j=i+1}^n \left\lceil \frac{R_i}{T_j} \right\rceil * C_j$$



Static (off-line) scheduling



Scheduling table:

■ Preemptive rate-monotonic scheduling

- Select tasks by increasing activation periods
- Schedule each job of the task for the hyperperiod ASAP
- **Exact** schedulability criterion: worst-case response-time analysis


- $j > i \iff T_j \leq T_i$

- Fixed-point computation (step 2):

- $$R_0^2 = C_0 + \text{ceil}(R_0^1/T_1) * C_1 + \text{ceil}(R_0^1/T_2) * C_2$$
$$= 3 + \text{ceil}(3/5) * 1 + \text{ceil}(3/5) * 2 = 6$$

$$R_i = C_i + \sum_{j=i+1}^n \left\lceil \frac{R_i}{T_j} \right\rceil * C_j$$

1 arrival of τ_1 and τ_2 in R_0^1



Static (off-line) scheduling



Scheduling table:

■ Preemptive rate-monotonic scheduling

- Select tasks by increasing activation periods
- Schedule each job of the task for the hyperperiod ASAP
- **Exact** schedulability criterion: worst-case response-time analysis


- $j > i \iff T_j \leq T_i$

- Fixed-point computation (step 3):

- $$R_0^3 = C_0 + \text{ceil}(R_0^2/T_1) * C_1 + \text{ceil}(R_0^2/T_2) * C_2$$
$$= 3 + \text{ceil}(6/5) * 1 + \text{ceil}(6/5) * 2 = 9$$

$$R_i = C_i + \sum_{j=i+1}^n \left\lceil \frac{R_i}{T_j} \right\rceil * C_j$$

2 arrivals
of τ_1 and τ_2
in R_0^2



Static (off-line) scheduling



Scheduling table:

■ Preemptive rate-monotonic scheduling

- Select tasks by increasing activation periods
- Schedule each job of the task for the hyperperiod ASAP
- **Exact** schedulability criterion: worst-case response-time analysis

- $j > i \iff T_j \leq T_i$

- Fixed-point computation (step 4):


- $R_0^4 = C_0 + \text{ceil}(R_0^3/T_1) * C_1 + \text{ceil}(R_0^3/T_2) * C_2$

- $= 3 + \text{ceil}(9/5) * 1 + \text{ceil}(9/5) * 2 = 9 = R_0^3$

- We reached a fixed point $\Rightarrow \text{WCRT}(\tau_0) = 9$

$$R_i = C_i + \sum_{j=i+1}^n \left\lceil \frac{R_i}{T_j} \right\rceil * C_j$$

2 arrivals
of τ_1 and τ_2
in R_0^3



Static (off-line) scheduling



Scheduling table:

- Preemptive rate-monotonic scheduling
 - Select tasks by increasing activation periods
 - Schedule each job of the task for the hyperperiod ASAP
 - **Exact** schedulability criterion: worst-case response-time analysis
 - Can be extended to account for blocking (critical sections in online schedulers), preemption costs, **multicore interference**

Static (off-line) scheduling



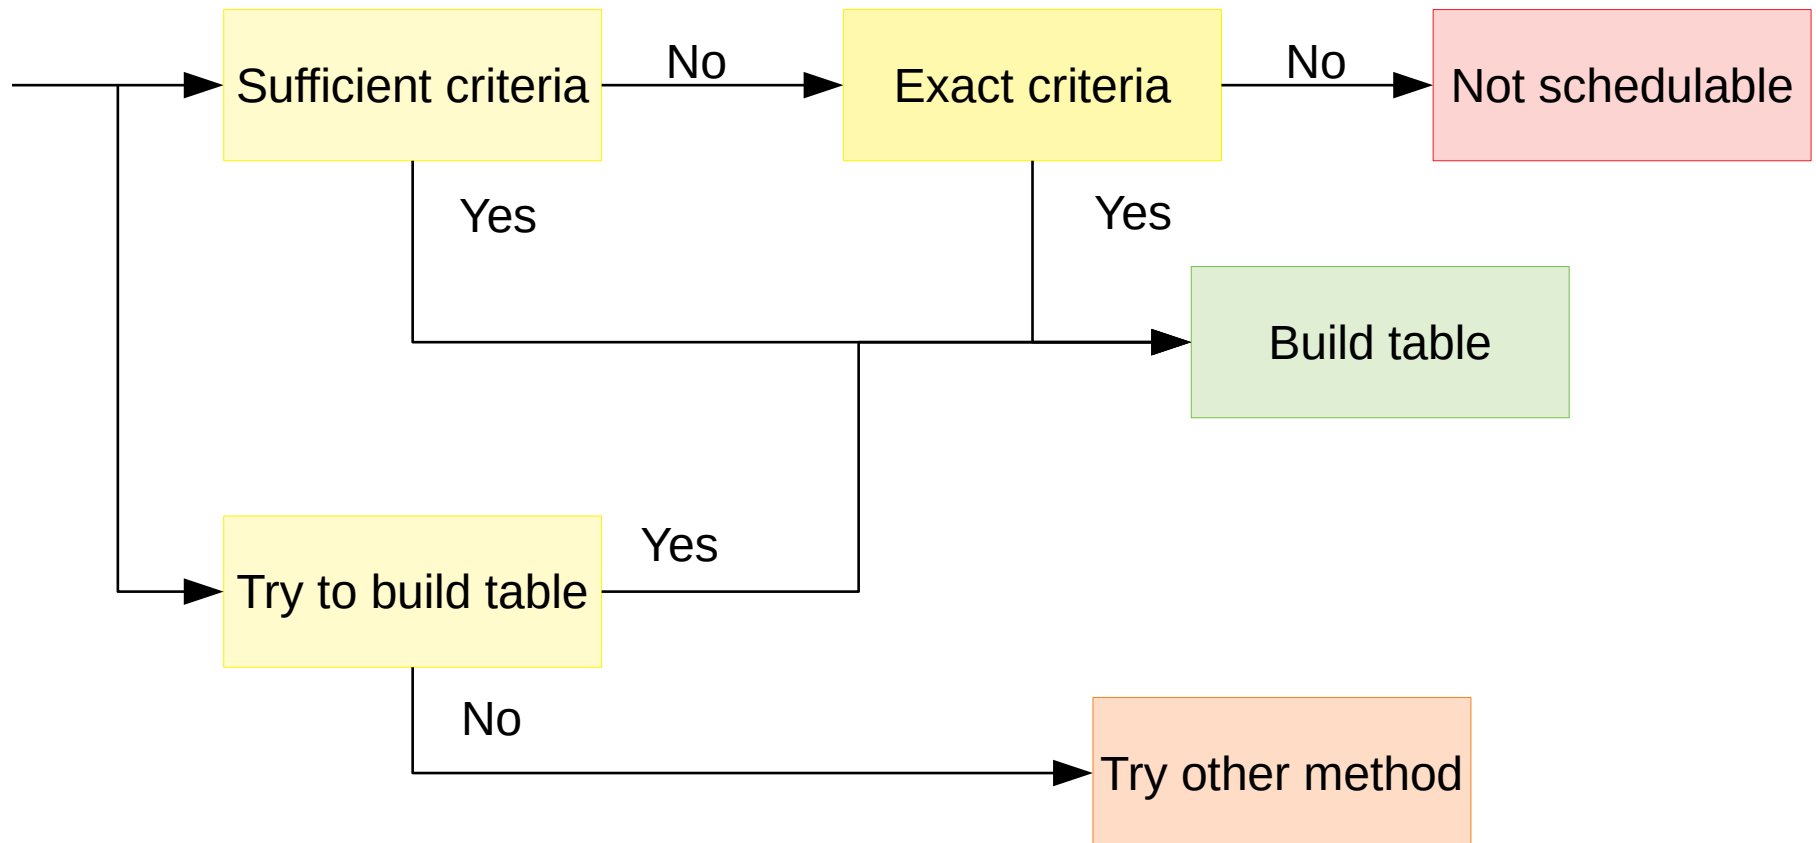
Scheduling table:

- Preemptive deadline-monotonic scheduling
 - Select tasks by increasing deadlines
 - Schedule each job of the task for the hyperperiod ASAP
 - When $D_i = T_i$, exactly the same as RM
 - Response-time analysis criterion can be obtained from the RM one :
 - Same formula
 - $j > i \iff D_j \leq D_i$ (priorities driven by deadlines)

Static (off-line) scheduling



Scheduling table: what do we do of this ?



Static (off-line) scheduling



Exercise 1:

■ Rate-monotonic scheduling

- Is the following task system RM-schedulable ?
 - $\tau_1 : C_1 = 1, T_1 = 10$
 - $\tau_2 : C_2 = 1, T_2 = 10$
 - $\tau_3 : C_3 = 2, T_3 = 20$
 - $\tau_4 : C_4 = 2, T_4 = 60$
 - $\tau_5 : C_5 = 2, T_5 = 60$
 - $\tau_6 : C_6 = 3, T_6 = 30$
 - $\tau_7 : C_7 = 4, T_7 = 60$
- If so, draw the scheduling table for the first hyperperiod of the system, using the RM algorithm.

Static (off-line) scheduling



Exercise 2:

■ Rate-monotonic scheduling

- Is the following task system RM-schedulable ?

- $\tau_1 : C_1 = 2, T_1 = 10$
- $\tau_2 : C_2 = 3, T_2 = 20$
- $\tau_3 : C_3 = 10, T_3 = 50$
- $\tau_4 : C_4 = 13, T_4 = 100$

Static (off-line) scheduling



Exercise 3:

- Rate-monotonic scheduling
 - Is the following task system RM-schedulable ?
 - $\tau_1 : C_1 = 12, T_1 = 50$
 - $\tau_2 : C_2 = 10, T_2 = 40$
 - $\tau_3 : C_3 = 10, T_3 = 30$

Static (off-line) scheduling



Exercise 4:

■ Rate-monotonic scheduling

- Is the following task system RM-schedulable ?

- $\tau_1 : C_1 = 1, T_1 = 4$
- $\tau_2 : C_2 = 1, T_2 = 5$
- $\tau_3 : C_3 = 1, T_3 = 6$
- $\tau_4 : C_4 = 2, T_4 = 7$

Static (off-line) scheduling



Exercise 5:

■ Rate-monotonic scheduling

- Is the following task system RM-schedulable ?
 - $\tau_1 : C_1 = 1, T_1 = 5$
 - $\tau_2 : C_2 = 2, T_2 = 10$
 - $\tau_3 : C_3 = 2, T_3 = 20$
 - $\tau_4 : C_4 = 5, T_4 = 10$
- If so, draw the scheduling table for the first hyperperiod of the system, using the RM algorithm.

Static (off-line) scheduling



Exercise 6:

- Rate-monotonic scheduling
 - Is the following task system RM-schedulable ?
 - $\tau_1 : C_1 = 3, T_1 = 7$
 - $\tau_2 : C_2 = 3, T_2 = 12$
 - $\tau_3 : C_3 = 5, T_3 = 20$

Static (off-line) scheduling



Exercise 7:

- Rate-monotonic scheduling
 - Is the following task system RM-schedulable ?
 - Process P : $C_1 = 1, T_1 = 3$
 - Process Q : $C_2 = 2, T_2 = 6$
 - Process S : $C_3 = 5, T_3 = 18$

Static (off-line) scheduling



Exercise 8:

- Consider three tasks P, Q and S
 - P : $C = 30\text{ms}$, $T = 100\text{ms}$
 - Q : $C = 1\text{ms}$, $T = 5\text{ms}$
 - S : $C = 5\text{ms}$, $T = 25\text{ms}$
- P is the most « important » task in the system, followed by Q and S
- What is the behavior of the scheduler if the priorities are based on the importance ?
- What is the combined utilisation U of tasks P, Q and S ?
- How should the processes be scheduled so that all deadlines are met ?
- Illustrate one of the schemes that allows these processes to be scheduled

Static (off-line) scheduling



Practical work:

- Develop a static scheduler
 - Input: multi-periodic task system described in text file
 - Policy: preemptive RM or DM
 - Output: textual description of the schedule
- Include a mechanism to guarantee the correctness of the produced schedule

Static (off-line) scheduling



Practical work:

- Develop a static scheduler
 - Input: multi-periodic task system described in text file **<= json format**
 - Policy: preemptive RM or DM
 - Output: textual description of the schedule **<= json format**
- Include a mechanism to guarantee the correctness of the produced schedule
- Use the python type Task to describe the RT tasks :

```
class Task():
```

```
    def __init__(self, name, idx, period, deadline, wcet):
```

```
        self.name=name
```

```
        self.idx=idx
```

```
        self.period=period
```

```
        self.deadline=deadline
```

```
        self.wcet=wcet
```

Static (off-line) scheduling



Practical work:

- Develop a static scheduler
 - Input: multi-periodic task system described in text file `<= json format`
 - Policy: preemptive RM or DM
 - Output: textual description of the schedule `<= json format`
- Include a mechanism to guarantee the correctness of the produced schedule
- Use the python type Task to describe the RT tasks
- Use the python function `parse_tasks()` to read a task system from a json file and obtain a python list of Tasks

Static (off-line) scheduling



Practical work:

- Develop a static scheduler
 - Input: multi-periodic task system described in text file `<= json format`
 - Policy: preemptive RM or DM
 - Output: textual description of the schedule `<= json format`
- Include a mechanism to guarantee the correctness of the produced schedule
- Use the python type Task to describe the RT tasks
- Use the python function `parse_tasks()` to read a task system from a json file and obtain a python list of Tasks
- Declare a python type to represent the schedule
- Write a method to generate a schedule from a list of tasks
- Write a method which verifies that a generated schedule is valid
- Write a method to export a schedule to a json file